INTERNAL SCORING RUBRIC - Data Scientist Engineer (Job ID: DJ-202)

SCORE SCALE

Scores are reported on a 0-100 scale with three criterion-referenced rating bands: strong_match (70-100), possible_match (40-69), and low_match (0-39). These bands align with the scoring pipeline's own thresholds, so the rubric and the scorer apply the same cut points.

MANDATORY REQUIREMENT

All mandatory conditions are required. A candidate missing any of them should not be rated above possible_match, and the score should not exceed 60.

- At least 3 years of commercial experience with hands-on Python development and NLP.
- Practical knowledge of statistics, machine learning and deep learning.
- Experience with a deep-learning framework such as PyTorch or TensorFlow.

PREFERRED CRITERIA

- Implementation of both classical NLP and neural-network-based NLP
- Familiarity with cloud ML platforms
- Upper-Intermediate or higher level of English

BASIS

Mandatory and preferred criteria are derived directly from the posted job requirements for this role.